



CS 165

Data Systems

Have fun learning to design and build modern data systems

class 8

indexing & sorting

prof. Stratos Idreos

[HTTP://DASLAB.SEAS.HARVARD.EDU/CLASSES/CS165/](http://daslab.seas.harvard.edu/classes/cs165/)



10/21

midterms

how to prepare

open book, notes, laptop/tablet
material from lectures, “browse/read” readings
check all quizzes and questions

quiz-like questions - no exact answer

10/21

midterms

how to prepare

open book, notes, laptop/tablet
material from lectures, “browse/read” readings
check all quizzes and questions

quiz-like questions - no exact answer

explain all steps and tradeoffs

expectations: describe the design space - chose what you think is the best approach (>1 if we ask for it) and then analyze in detail all requests - if you made the wrong choice in the beginning it is OK - but say so if you find out in the end and explain as much as possible

10/21

midterms

how to prepare

open book, notes, laptop/tablet
material from lectures, “browse/read” readings
check all quizzes and questions

quiz-like questions - no exact answer

explain all steps and tradeoffs

expectations: describe the design space - chose what you think is the best approach (>1 if we ask for it) and then analyze in detail all requests - if you made the wrong choice in the beginning it is OK - but say so if you find out in the end and explain as much as possible

we count best of Midterm 2 or Midterm1+2

vectorwised processing: how to


select max(A) from R where B<20

```
p=select(B,null,20)
a=fetch(A,p)
res=max(a)
```

vectorwised processing: how to

select max(A) from R where B<20

```
p=select(B,null,20)
a=fetch(A,p)
res=max(a)
```



rewrite to


```
j=0;
for(i=0; i<B.size; i+=vector.size){
    p=select(B,i,vector.size,null,20)
    a=fetch(A,p)
    rv[j++]=max(a)
}
res=max(rv)
```

edge cases
not included :)

vectorwised processing: how to

select max(A) from R where B < 20

```
p=select(B,null,20)
a=fetch(A,p)
res=max(a)
```



```
j=0;
for(i=0; i<B.size; i+=vector.size){
    p=select(B,i,vector.size,null,20)
    a=fetch(A,p)
    rv[j++]=max(a)
}
res=max(rv)
```

edge cases
not included :)

take plans
from here:

Extra: **Enhanced stream processing in a DBMS kernel**

Erietta Liarou, Stratos Idreos, Stefan Manegold, Martin Kersten

In Proc. of the International Conf. on Extending Database Technology, 2013

essential column-stores features

virtual ids

late tuple reconstruction (if ever)

vectorized execution

compression

fixed-width columns

essential column-stores features


virtual ids

late tuple reconstruction (if ever)

vectorized execution

compression

fixed-width columns



Column-stores vs. row-stores: how different are they really?

D. Abadi, S. Madden, and N. Hachem


ACM SIGMOD Conference on Management of Data, 2008

but why now...

weren't all those design options obvious in the past as well?


but why now...

weren't all those design options obvious in the past as well?



but why now...

weren't all those design options obvious in the past as well?



- 1) big memories**
- 2) cpu vs memory speed**


main-memory systems

optimized for the memory wall
with or without persistent data

IBM

ORACLE®

SAP



Microsoft

cloudera®
IMPALA

APACHE
Spark™


Hadoop

 **Qaction**®

twitter



facebook.

 snowflake

Google™



Linked**in**™

 **platfora**®

mem**sql**

 **hp**  **VERTICA**




HARVARD
School of Engineering
and Applied Sciences

CS165, Fall 2019
Stratos Idreos

IBM

ORACLE®

SAP



Microsoft

cloudera®
IMPALA

APACHE
Spark™



Qaction®

twitter



facebook.

snowflake

Google™




LinkedIn™

platfora®

memsql



column-stores = bad name

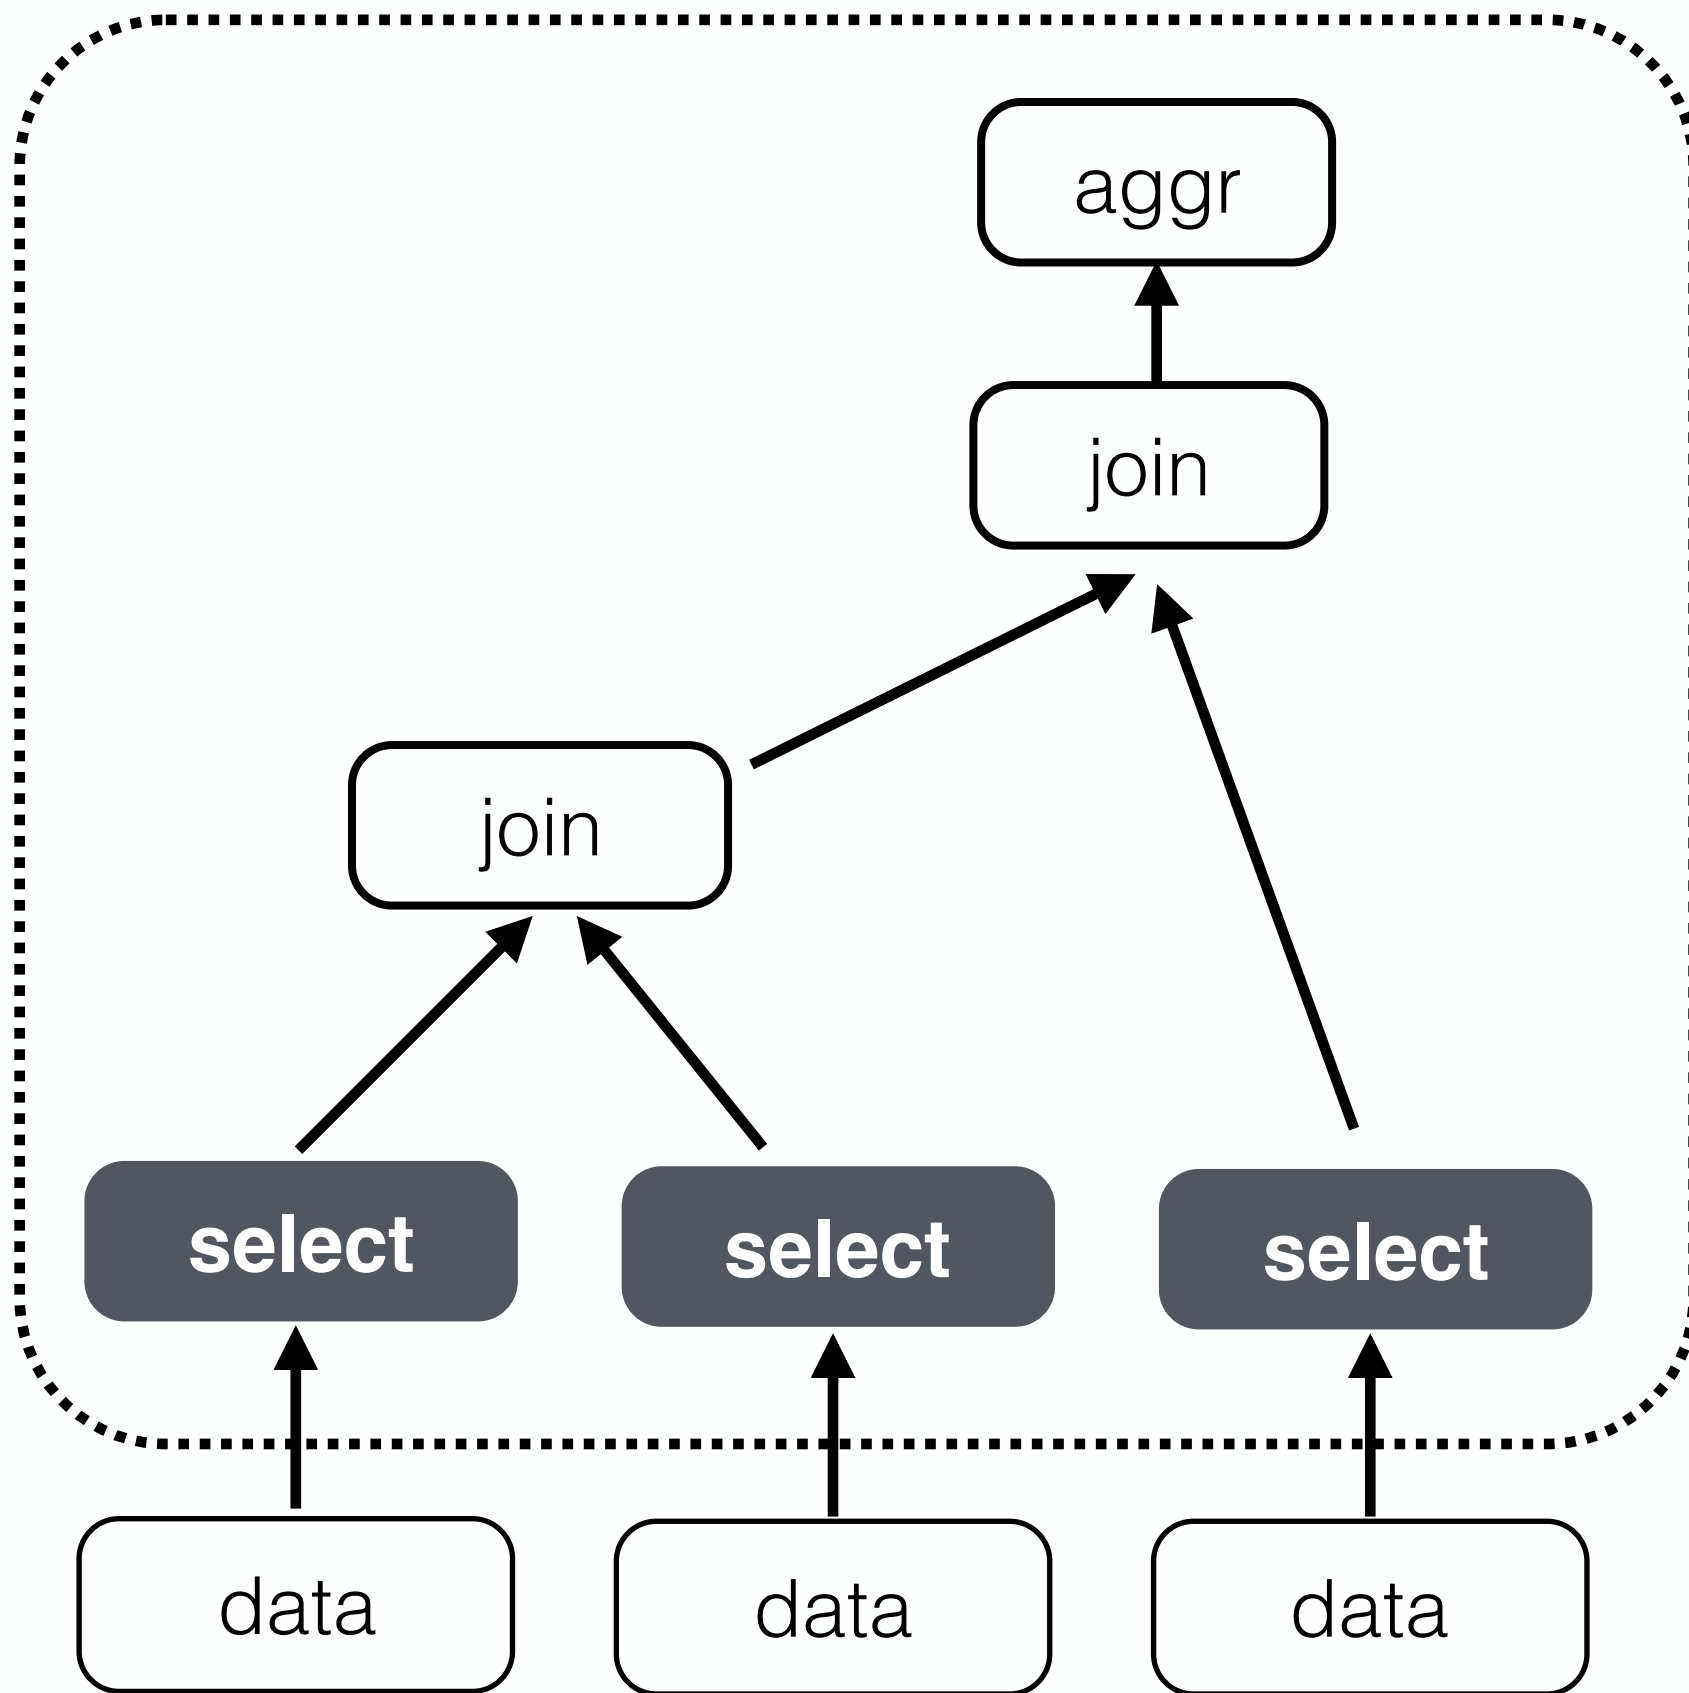


analytical systems



first part done: basic concepts in modern systems


coming up: indexing and fast scans



**it all starts with the
select operator**

it touches all the data






index knows structure of the data
filtering data: point/range queries

an alternative data representation


why not just sort the data?




index knows structure of the data
filtering data: point/range queries

an alternative data representation

let's go with sorting for a while






initial state
columns in
insertion order






select B+C from R
where $A < 10$


values are out of order





values are out of order



intermediate out of order




initial state
columns in
insertion order



**propagate
order of A**



select max(D),min(E) **from** R **where** (A>10 and A<40) and (B>20 and B<60)




binary search
for 10 & 40


for all B values between
pos1 & 2: if B>20 and B<60
mark bit vector at pos i

for each marked
position
max(D)

select max(D),min(E) **from** R **where** (A>10 and A<40) and (B>20 and B<60)





select max(D),min(E) **from** R **where** (A>10 and A<40) **or** (B>20 and B<60)




binary search
for 10 & 40

for all B values **outside**
pos1 & 2: if B>20 and B<60
mark bit vector at pos i

for each marked
position
max(D)



queries that filter
on A benefit



queries that filter
on B benefit

...

C-Store: A Column-oriented DBMS

Michael Stonebraker, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, Amerson Lin, Samuel Madden, Elizabeth J. O'Neil, Patrick E. O'Neil, Alex Rasin, Nga Tran, Stanley B. Zdonik

In Proc. of the Very Large Databases Conference (**VLDB**), 2005

A B C

A B C

queries that filter

Column-store Projections

We can have **many** of them to fit different access patterns

A B C

...

C-Store: A Column-oriented DBMS

Michael Stonebraker, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, Amerson Lin, Samuel Madden, Elizabeth J. O'Neil, Patrick E. O'Neil, Alex Rasin, Nga Tran, Stanley B. Zdonik

In Proc. of the Very Large Databases Conference (**VLDB**), 2005

A B C

A B C

queries that filter

Column-store Projections

We can have **many** of them to fit different access patterns

But there are many possible ones...how to choose?

A B C


...

C-Store: A Column-oriented DBMS


Michael Stonebraker, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, Amerson Lin, Samuel Madden, Elizabeth J. O'Neil, Patrick E. O'Neil, Alex Rasin, Nga Tran, Stanley B. Zdonik

In Proc. of the Very Large Databases Conference (**VLDB**), 2005

declarative interface
ask what you want



declarative interface
ask what you want



**declarative
paradigm
is broken**

online ? 

storage budget \ll smaller than the
possible set of projections



storage budget \ll smaller than the
possible set of projections

incrementally, adaptively create partial projections

Browse: **Self-organizing tuple reconstruction in column-stores**

Stratos Idreos, Martin Kersten, Stefan Manegold

In Proc. of the ACM **SIGMOD** Inter. Conference on Management of Data, 2009



Sorting is used to create & maintain projections but also across numerous other operations in a system. How can we sort efficiently over large data and modern hardware? Assume an array of N integers and a two level memory hierarchy.

CPU

L1 memory

L2 memory

cost to sort an array C_s ?

cost to find a value once sorted C_a ?


optimized algorithm to minimize C_s & C_a

data does not fit in L1 memory; it fits in L2
CPU can read/write directly from/to L1 only

memory
level L

(size=3 pages)


memory
level L+1



initial state: 8 unordered pages

quicksort in place

memory
level L



(size=3 pages)

memory
level L+1




initial state: 8 unordered pages

memory
level L

(size=3 pages)


memory
level L+1



initial state: 8 unordered pages


quicksort in place

memory
level L



(size=3 pages)

memory
level L+1




initial state: 8 unordered pages

memory
level L

(size=3 pages)

memory
level L+1



initial state: 8 unordered pages


quicksort in place



memory
level L

(size=3 pages)

memory
level L+1




initial state: 8 unordered pages

each page is now sorted
we read and wrote every page once
data movement cost is $2N$ pages

memory
level L

(size=3 pages)

memory
level L+1




initial state: 8 unordered pages

memory
level L


(size=3 pages)

memory
level L+1



initial state: 8 sorted pages


merge to new page



memory
level L


(size=3 pages)

memory
level L+1



initial state: 8 sorted pages


merge to new page



memory
level L


(size=3 pages)

memory
level L+1



initial state: 8 sorted pages


merge to new page



memory
level L

(size=3 pages)

memory
level L+1



initial state: 8 sorted pages


merge to new page



memory
level L


(size=3 pages)

memory
level L+1



initial state: 8 sorted pages


merge to new page



memory
level L

(size=3 pages)

memory
level L+1




initial state: 8 sorted pages

memory
level L

(size=3 pages)

memory
level L+1




initial state: 8 sorted pages

each pair of pages is now sorted
we read and wrote every page once
data movement cost is $2N$ pages (total $2N+2N$)

memory
level L

(size=3 pages)

memory
level L+1



initial state: 8 sorted pages



1 pass to merge into **8** sorted pages ($2N$ pages)

1 pass to merge into **4** sorted pages ($2N$ pages)

1 pass to merge into **2** sorted pages ($2N$ pages)


1 pass to sort each page ($2N$ pages)

1 pass to merge into 8 sorted pages (2N pages)

1 pass to merge into 4 sorted pages (2N pages)

1 pass to merge into 2 sorted pages (2N pages)

1 pass to sort each page (2N pages)





1 pass to merge into 8 sorted pages (2N pages)

1 pass to merge into 4 sorted pages (2N pages)

1 pass to merge into 2 sorted pages (2N pages)

1 pass to sort each page (2N pages)



$$\log_2(N)+1$$

1 pass to merge into 8 sorted pages (2N pages)

1 pass to merge into 4 sorted pages (2N pages)

1 pass to merge into 2 sorted pages (2N pages)

1 pass to sort each page (2N pages)



$$2N(\log_2(N)+1)$$



1 pass to merge into 8 sorted pages (2N pages)

1 pass to merge into 4 sorted pages (2N pages)

1 pass to merge into 2 sorted pages (2N pages)

1 pass to sort each page (2N pages)



$$2N(\log_2(N)+1) \times \text{bytesPerPage}$$



we have M pages in memory (not just 3)

we have M pages in memory (not just 3)

$$2N(\log_2(N)+1)$$

we have M pages in memory (not just 3)

$$2N(\log_2(N)+1) \rightarrow 2N(\log_{M-1}(N)+1)$$

we have M pages in memory (not just 3)

$$2N(\log_2(N)+1) \rightarrow 2N(\log_{M-1}(N)+1)$$

immediately sort groups of M pages in first pass

we have M pages in memory (not just 3)

$$2N(\log_2(N)+1) \rightarrow 2N(\log_{M-1}(N)+1)$$

immediately sort groups of M pages in first pass

$$2N(\log_{M-1}(N)+1)$$

we have M pages in memory (not just 3)

$$2N(\log_2(N)+1) \rightarrow 2N(\log_{M-1}(N)+1)$$

immediately sort groups of M pages in first pass

$$2N(\log_{M-1}(N)+1) \rightarrow 2N(\log_{M-1}(N/M)+1)$$

previous discussion holds for all
levels of memory hierarchy

other usage of sorting, e.g.,:


order by

group by

sort-merge join

remove duplicates

sort/cluster ids/positions to avoid random access



data size: N pages
memory size: M pages

how much memory M do we
need to sort N data in p passes only?

or

how much data can we sort in
 p passes if we have M memory?

$$\log_{M-1}(N/M) + 1 \leq p$$



Read **textbook**: Chapter 13

Browse: **Self-organizing tuple reconstruction in column-stores**

Stratos Idreos, Martin Kersten, Stefan Manegold

In Proc. of the ACM **SIGMOD** Inter. Conference on Management of Data, 2009

indexing & sorting

DATA SYSTEMS

prof. Stratos Idreos

