

Recitation for indexing

Part1 Btree

Storage Types

→ Cache

- ◆ Fastest/most costly; volatile;

→ Main Memory

- ◆ Fast access; too small for entire db; volatile

→ Disk

- ◆ Long-term storage of data; random access; non-volatile

→ Flash Memory

- ◆ No seeks (cheap reads);

How Disk Works

- Surface divided into tracks -> sectors (smallest unit of data that can be R/W)
 - ◆ Disk arm swings to position head on right track
 - ◆ Platter spins continually as data is R/W from sector
- Measuring Disk Speed
 - ◆ Access Time:
 - Seek Time: time to find right track
 - Latency Time: time for sector to appear under head
 - ◆ Data Transfer Rate:
 - The rate at which data can be retrieved
- Sequential I/O < Random I/O

Problem Statement

- Many queries reference only a fraction of records
- It is inefficient for the system to read every single tuple
- A better design is for the system to locate reference directly

Index

→ Indexes


- ◆ Auxiliary data structures over relations that can improve the search time
- ◆ Think about index in a textbook
- ◆ We search for an index, find corresponding pages, and then read information to find what we are looking for
- ◆ The index is much smaller than the book and has words sorted in alphabetical order

B-trees

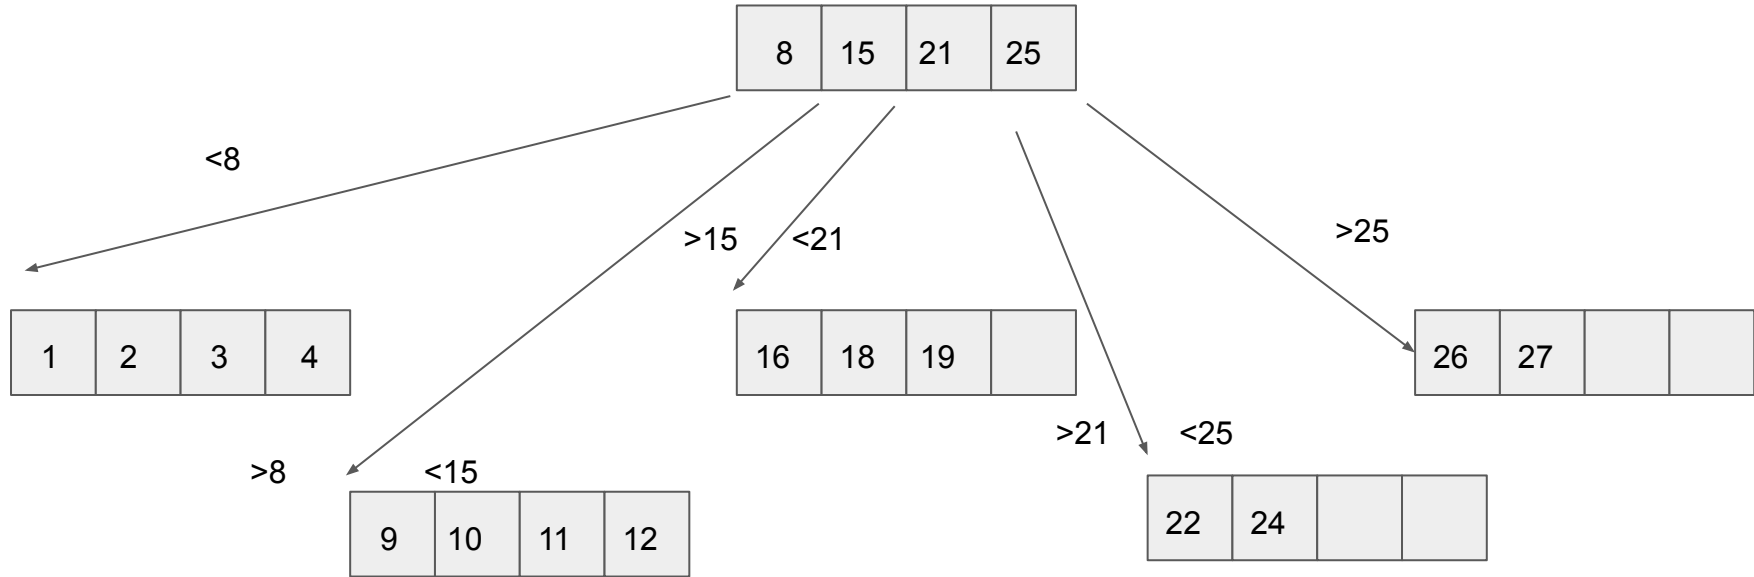
- Most successful family of index schemes (B-trees, B+-trees,...)
- Can be used for primary/secondary, clustering/non-clustering index
- Balanced “n-way” search trees
 - ◆ Helpful to think of it as a tree structure with n-pointers in each node

B-tree nodes

- Key values are ordered
- MAXIMUM: n pointers
- MINIMUM: $\lceil n/2 \rceil$ pointers
 - ◆ Exception: root's minimum = 2
- Internal nodes must have $n-1$ key values where n is the number of pointers



B-tree of order 5

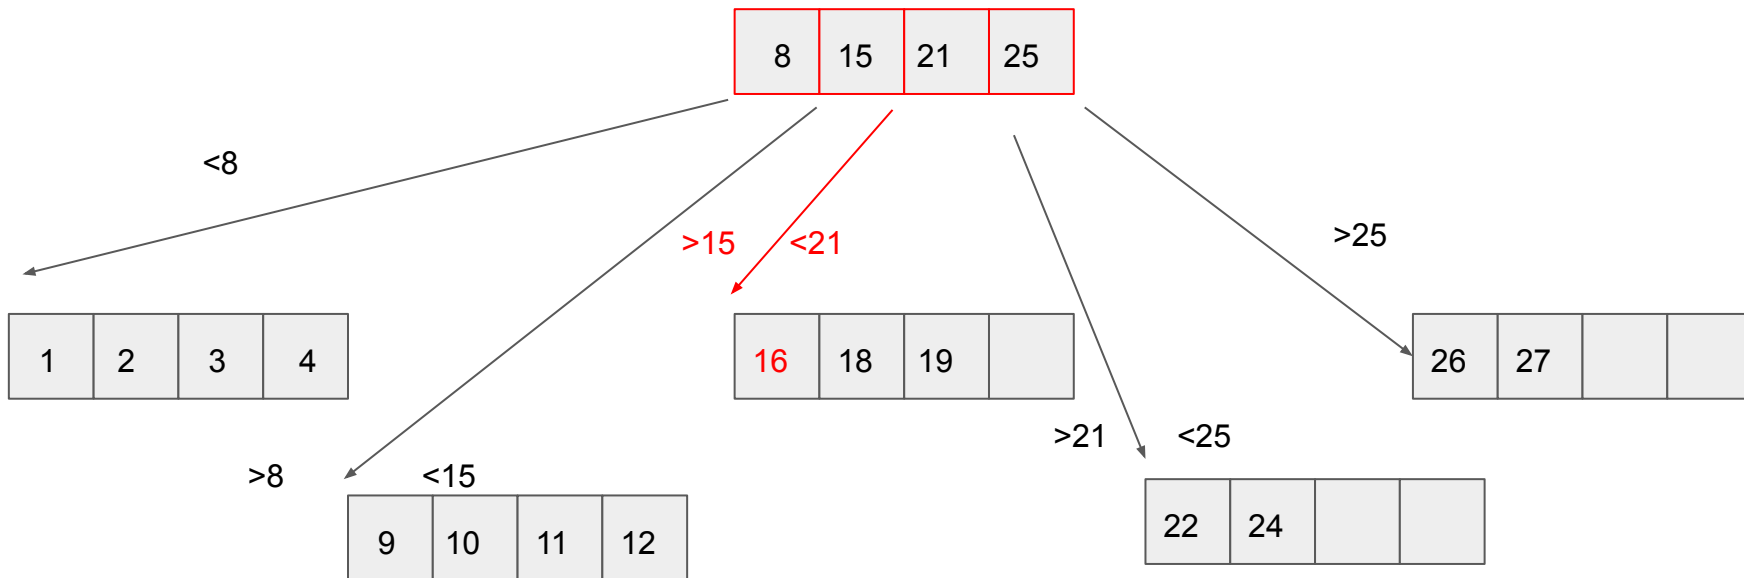


- Key values appear once
- Record pointers accompany keys

Queries

→ Let's run through an exact match query

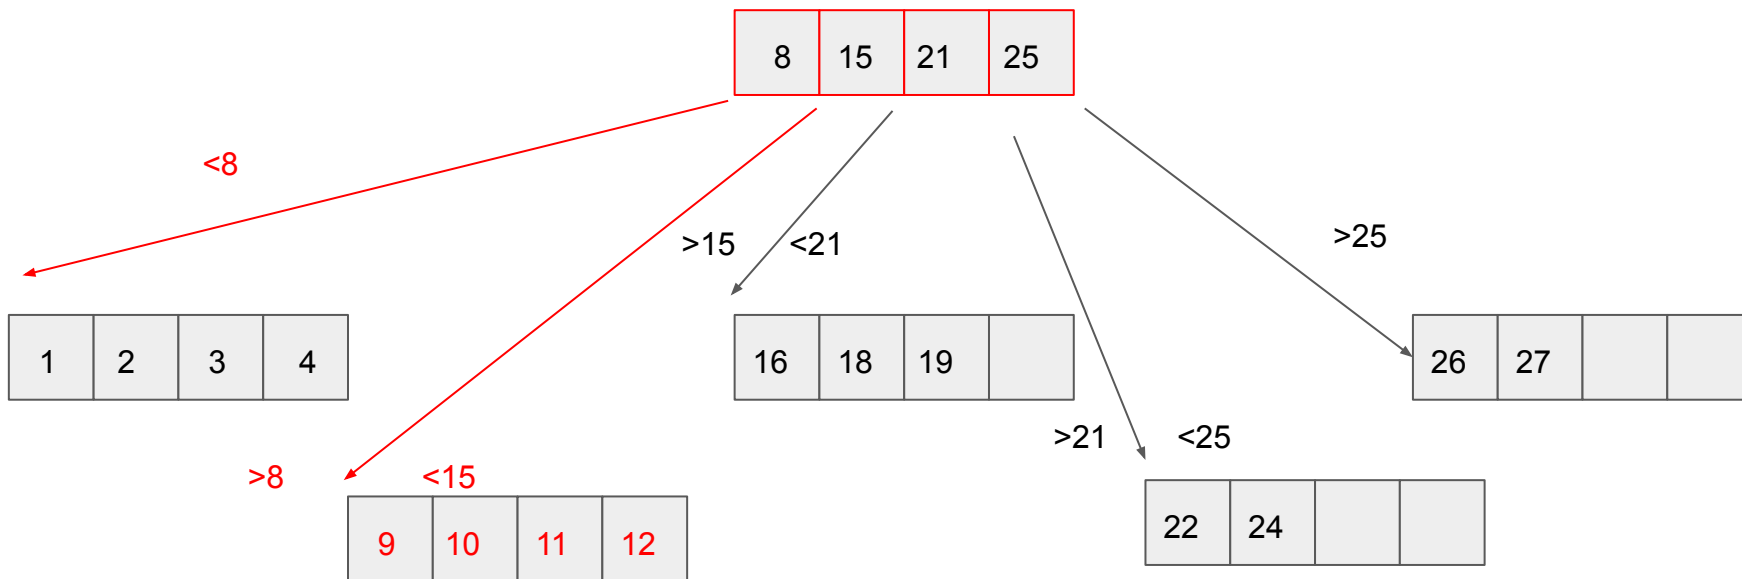
◆ Let's find the value 16?



Queries

→ Let's run through a range query

◆ Let's find the values between 6 and 13?



How To Maintain B-trees?

→ B-tree rules must be obeyed on every insert + delete

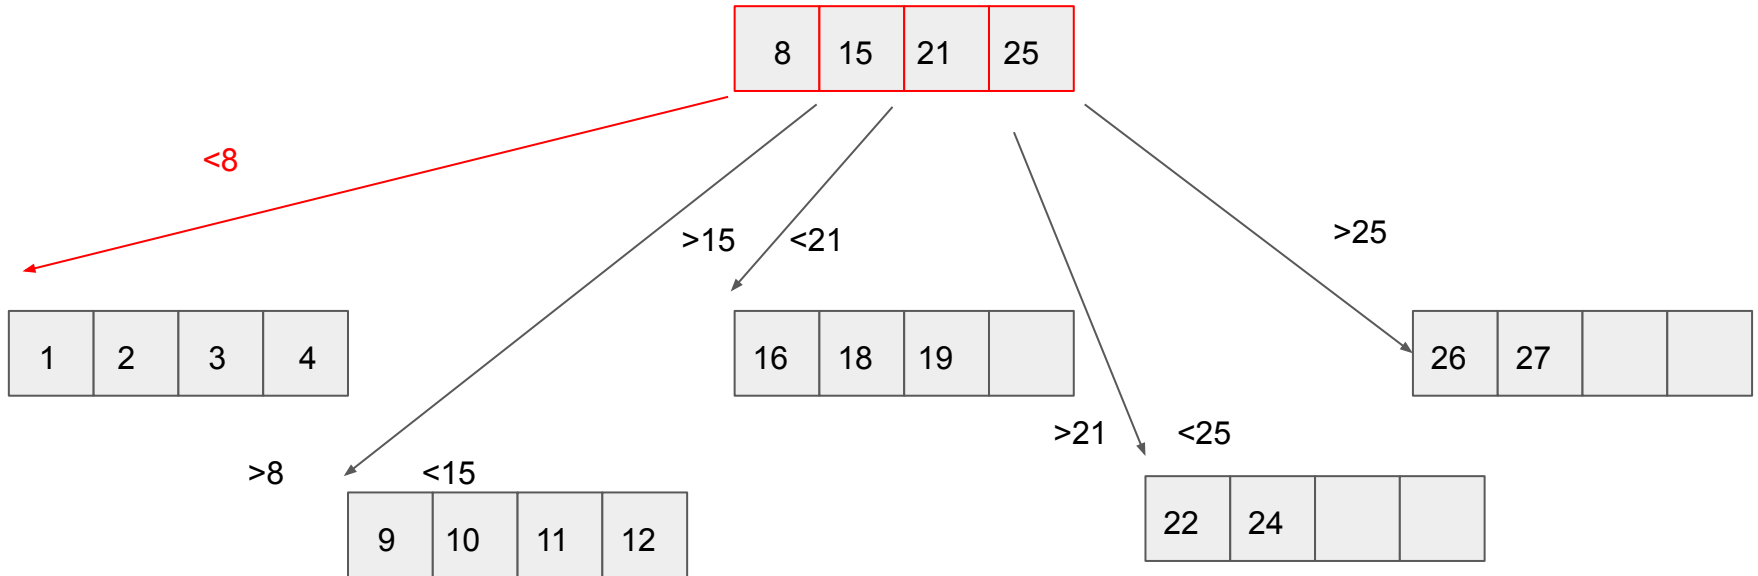
→ Rules:

- ◆ Insert in leaf, if room exists
- ◆ On overflow (no more room)
 - Split: create a new internal node
 - Redistribute keys
 - So that it preserves B-tree properties
 - Push middle key up (recursively)

Queries

→ Let's run through an overflow insertion

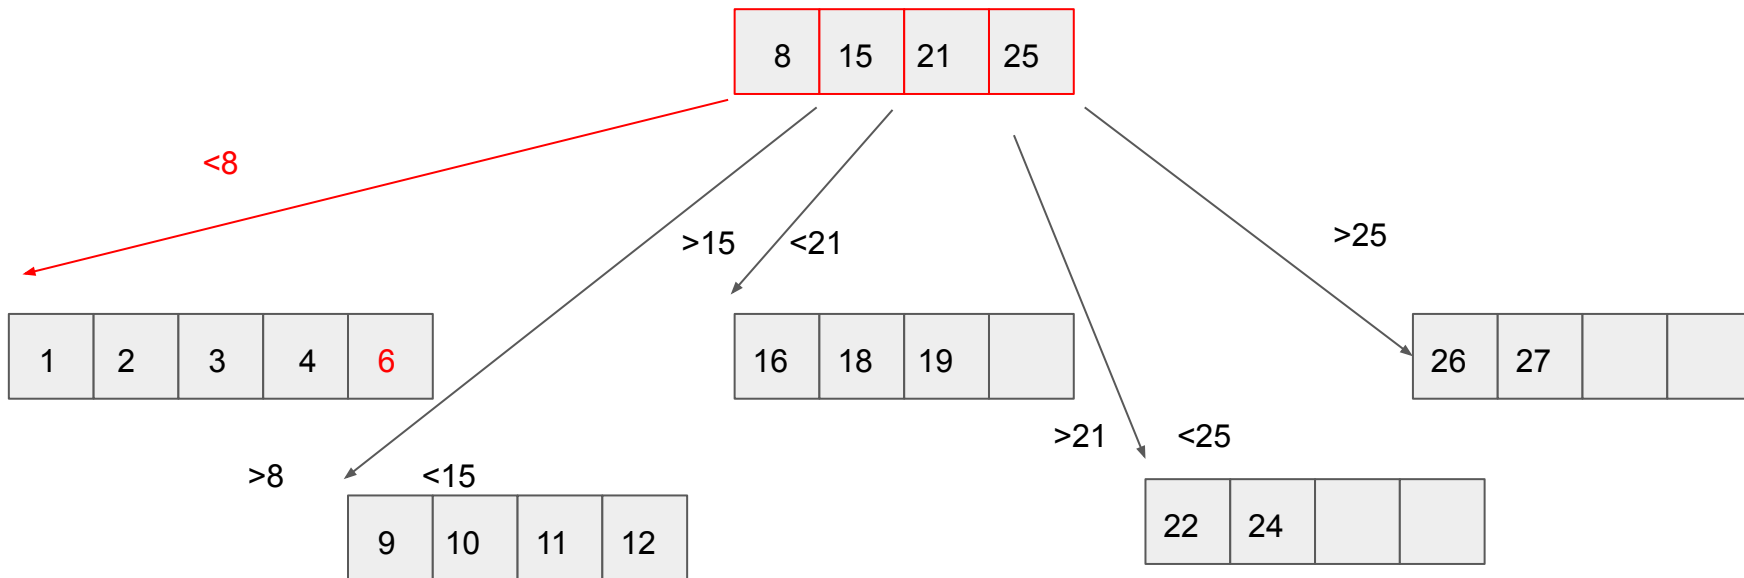
◆ Let's insert the value 6:



Queries

→ Let's run through an overflow insertion

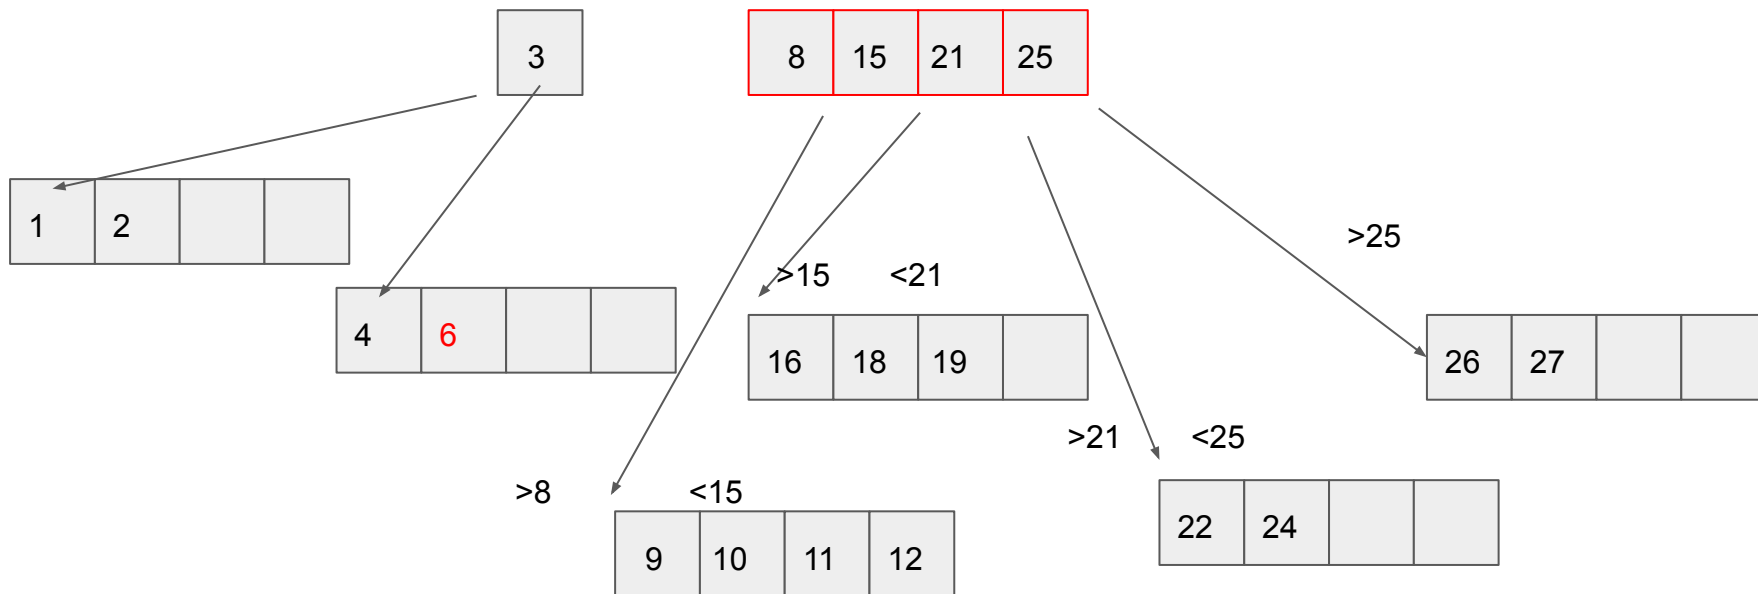
◆ Let's insert the value 6:



Queries

→ Let's run through an overflow insertion

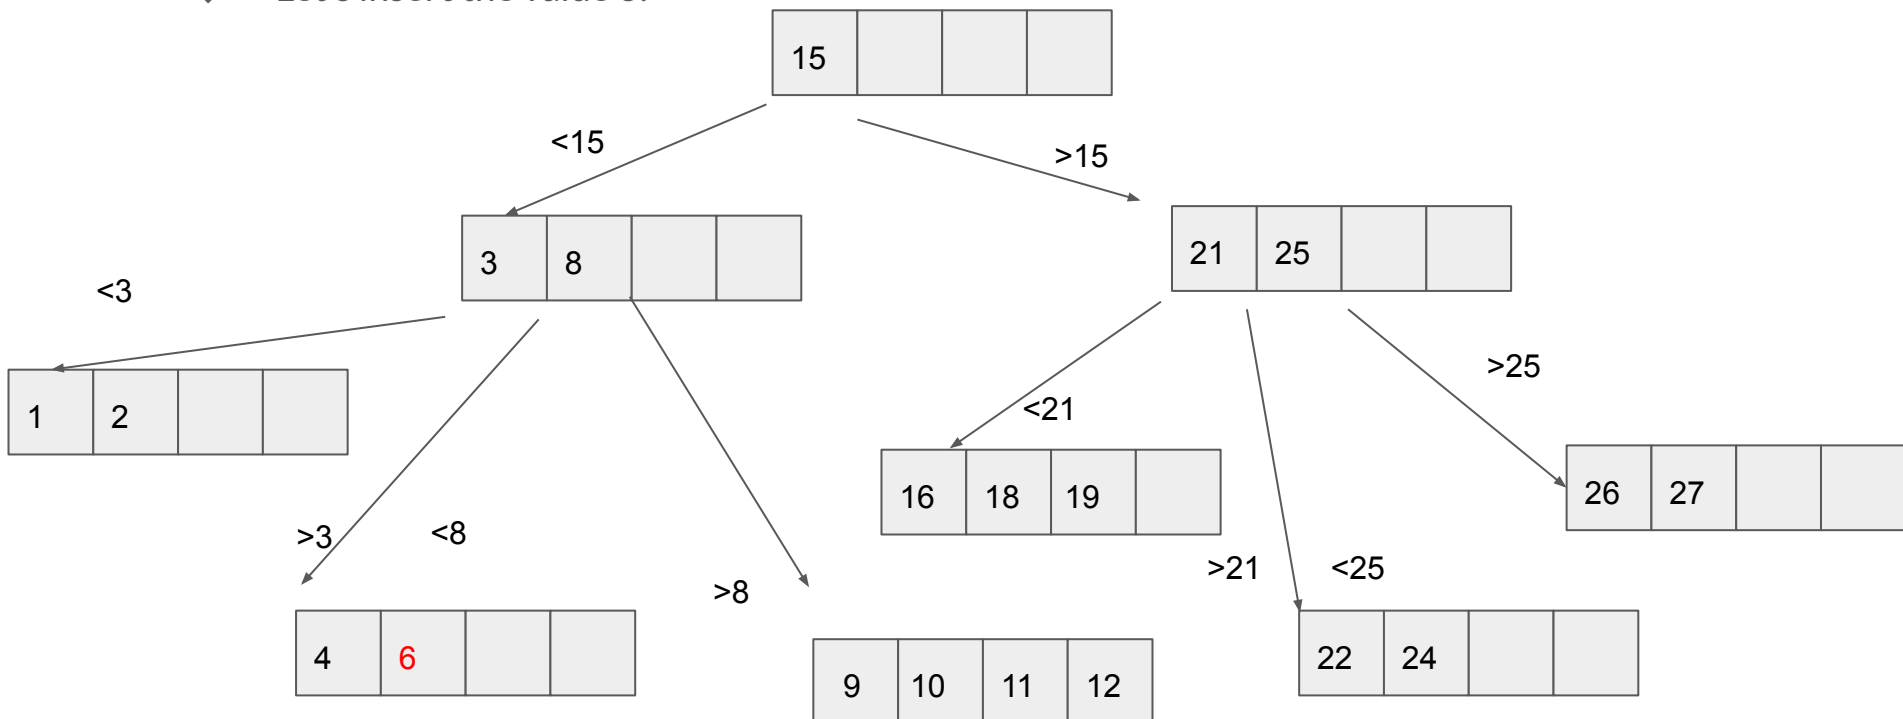
◆ Let's insert the value 6:



Queries

→ Let's run through an overflow insertion

◆ Let's insert the value 6:




B+-trees

- Allow sequential operations
 - ◆ String all leaf nodes together
 - ◆ Every key appears at leaf level (some keys show up more than once)
- This means non-leaf nodes do not contain data for the key, only the pointer value and key value


B+-tree Insertion

→ Insert 5 into B-tree of order 3




B+-tree Insertion

→ Insert 5




B+-tree Insertion

→ Insert 5




B+-tree Insertion

→ Insert 5



B+-tree Insertion

→ Insert 5



Recitation for indexing

Part-2 hashing


When do we need hashing?

Hash-based indexes are best for **equality selections**. Cannot support range searches.


Two kinds of hashing:

- Static hashing: will degrade performance with long overflow chains
- Dynamic hashing: Extendible hashing and linear hashing

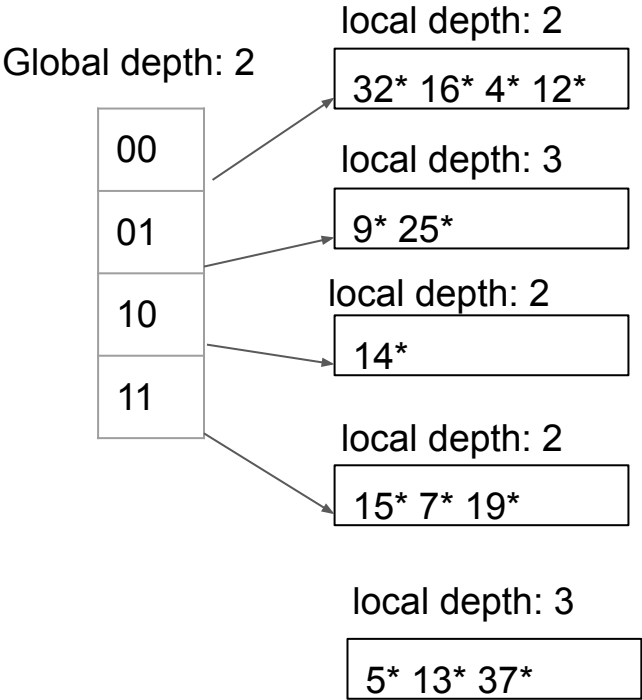
Extendible hashing



Extendible hashing


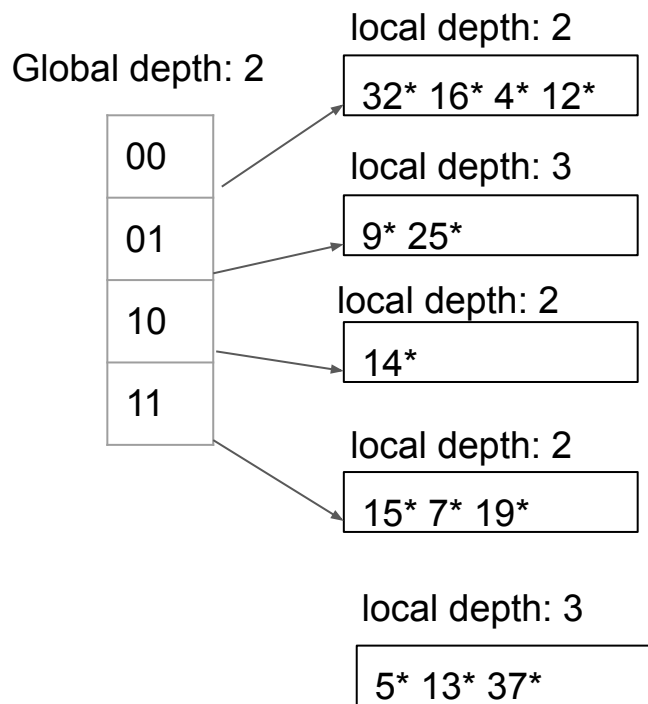


Insert 37




Extendible hashing

Insert 37



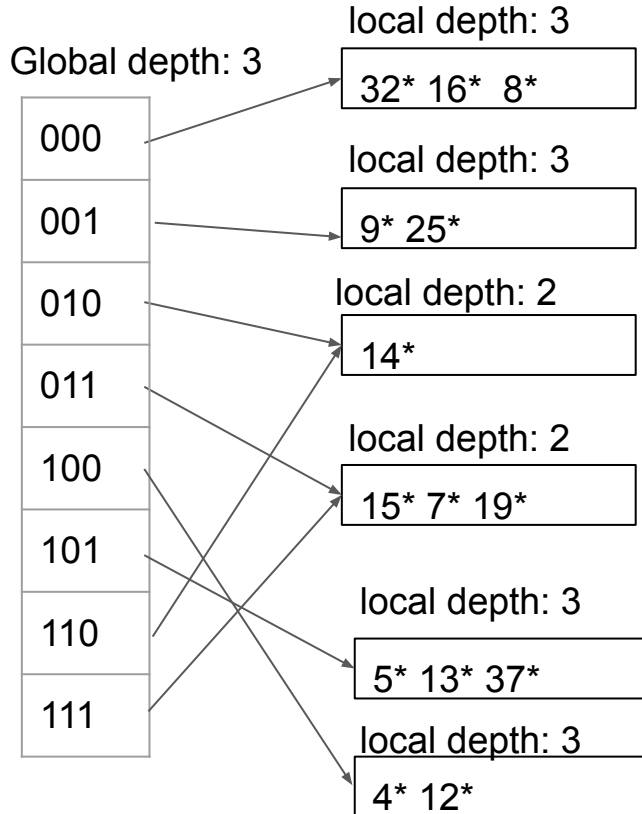
Extendible hashing

What will happen when Inserting 8*?



Extendible hashing


What will happen when Inserting 8*?




Linear Hashing

Insert 37

Level=0, next=0, n=4

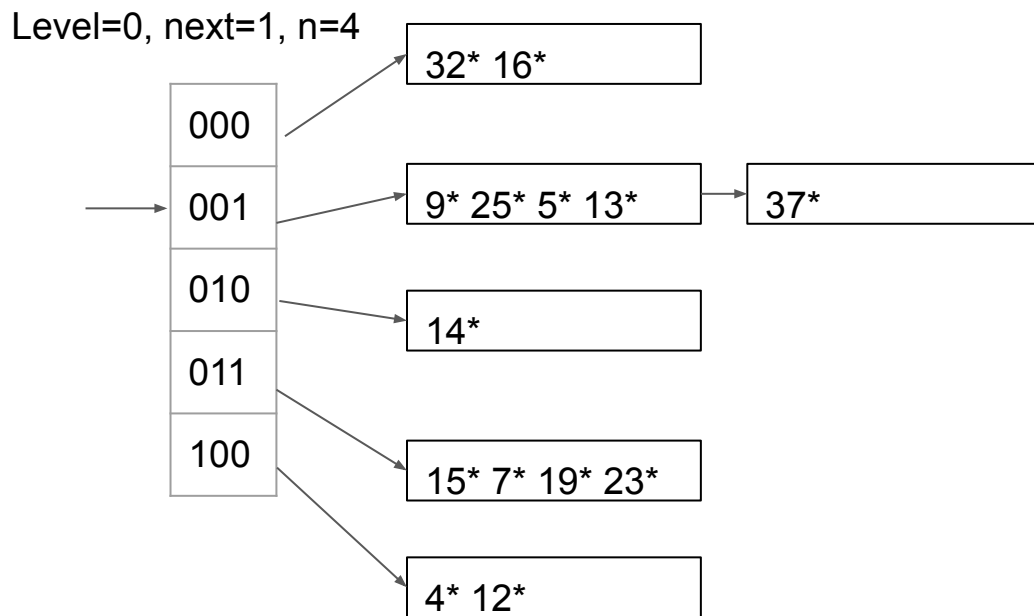


Level=0, next=1, n=4



Linear Hashing

What will happen when inserting 27?



Linear Hashing

What will happen when inserting 27?

