

# **CSCI1270**

# **Introduction to Database Systems**

with thanks to Prof. George Kollios, Boston University  
Prof. Mitch Cherniack, Brandeis University  
Prof. Avi Silberschatz, Yale University

# What is a Database System?

- Database:

A very large collection of *related* data

- Models a real world enterprise:

  - ↗ Entities

    - Sports: Teams, players, games
    - University: Students, professors, books, courses

  - ↗ Relationships

    - The Patriots are *playing* in the Superbowl

- DBMS (**D**atabase **M**anagement **S**ystem): A software system that can be used to store, manage, retrieve and transform data from a database.

- Database System: DBMS+data (+ applications)

# Why Study Databases?

## ■ Shift from computation to information

↗ Always true for corporate computing

↗ Scientific datasets:

- Astronomy

- Biology

- Particle Physics

↗ The Web

## ■ DBMS encompasses much of CS in a practical discipline

Operating systems

Languages

Distributed systems

Performance

Theory

AI

## ■ JOBS!

*High-paying jobs!!!*

# Why Databases?

## ■ Why not store everything in flat files?

i.e., use the file system of the OS: cheap, simple...

[ ***Name, Course, Grade*** ] ← File system does not know even this.

John Smith, CS22, B;

Mike Stonebraker, CS123, A;

Jim Gray, CS127, A;

John Smith, CS227, B+;

.....

This is how things were  
in the “Bad Old Days”

## ■ Simple? Yes, but not scalable...

Have to **scan** and **reparse** the file on every access.


# Scalability

- While you might build a database to store your favorite recipes, this is not the interesting case.
- Modern DBMS's are designed to scale  
(i.e., >> a terabyte)

Must be well optimized on a single node

Must be able to execute on 100's or 1000's of nodes.

# DBMS Workloads



Data is:

**mutable** => changes often

(1) data and (2) structure

**shared** between different users and applications

**persistent and long-lived** => must withstand failures  
evolvable

# Storage Problems

## Apps Own Files

### ■ Data redundancy and inconsistency

- Multiple file formats,
- duplication of information in different files

*MyApp ( **Name, Course, Email, Grade** )*

John Smith, CS112, [js@cs.bu.edu](mailto:js@cs.bu.edu), B

Jim Gray, CS560, [jg@cs.bu.edu](mailto:jg@cs.bu.edu), A

John Smith, CS560, [js@cs.bu.edu](mailto:js@cs.bu.edu), B+

*YourApp ( **Name, Email, Course, Grade** )*

Mike Stonebraker, [ms@cs.bu.edu](mailto:ms@cs.bu.edu), CS234, A

J. Smith, [js@cs.bu.edu](mailto:js@cs.bu.edu), CS560, B+

Why is this a problem?

- Wasted space (?)
- Potential inconsistencies

(e.g., multiple conventions, John Smith vs Smith J.)

# Retrieval Problems

## ■ Data retrieval (Query):

- ↗ Find the students who took CS560
- ↗ Find the students with  $\text{GPA} > 3.5$

For every query we would need to write a program!  
Each program would read the whole file.

## ■ We need the retrieval to be:

- ↗ Easy to write
- ↗ Execute efficiently



# Data Integrity

## ■ Data Integrity in flat file model

- ✚ Poor support for **sharing**:
  - Prevent simultaneous modifications
- ✚ Poor coping mechanisms for **system crashes**
- ✚ No means of Preventing **Data Entry Errors** (checks must be hard-coded in the programs)
- ✚ **Security** problems

## ■ Database systems offer solutions to all of the above problems

# Evolution of data

- Long-lived data → Evolution
- What happens if I need to change how the data is stored?
  - ↗ Access patterns change
  - ↗ Tuning
- Should not have to re-write all my applications.
- Solution: Data independence!

# Data Organization

- **Data Models:** a framework for describing
  - ↗ data objects
  - ↗ data relationships
  - ↗ data semantics
  - ↗ data constraints
- Presents primitives for
  - ↗ Representing Data → Data Definition Language (DDL)
  - ↗ Manipulating Data → Data Manipulation Language (DML)
- We will concentrate on the **relational model**
- Other models:
  - ↗ ***Entity-Relationship model*** (we will discuss)
  - ↗ object-oriented model
  - ↗ semi-structured data models, XML
  - ↗ Array data model

# Database Schema

- Similar to types and variables in programming languages
- **Schema** – the structure of the database
  - ↗ e.g., the database consists of information about a set of customers and accounts and the relationship between them
  - ↗ Expressed in some *data model* (DDL)
  - ↗ Occurs at multiple levels
    - **Logical schema**: database design at the logical level
    - **Physical schema**: database design at the physical level

# Levels of Data Modeling

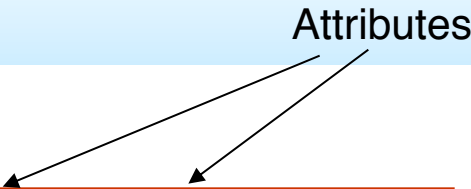
- **Logical level:** describes data stored in database, and the relationships among the data.

```
type customer = record
    name : string;
    street : string;
    city : integer;
end;
```

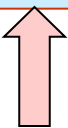
- **Physical level:** describes how a record is stored.
  - ↗ Sorting, page-alignment, indexes
- Also, **View level:** application programs hide details of data types. Views can also hide information (e.g., salary) for security purposes.

# Relational Model

- Example of tabular data in the relational model



<i>Customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>	<i>account-number</i>
192-83-7465	Johnson	Alma	Palo Alto	A-101
019-28-3746	Smith	North	Rye	A-215
192-83-7465	Johnson	Alma	Palo Alto	A-201
321-12-3123	Jones	Main	Harrison	A-217
019-28-3746	Smith	North	Rye	A-201



**Key**

**Schema** = **Customer** (Customer-id, customer-name, customer-street, customer-city, account-number)

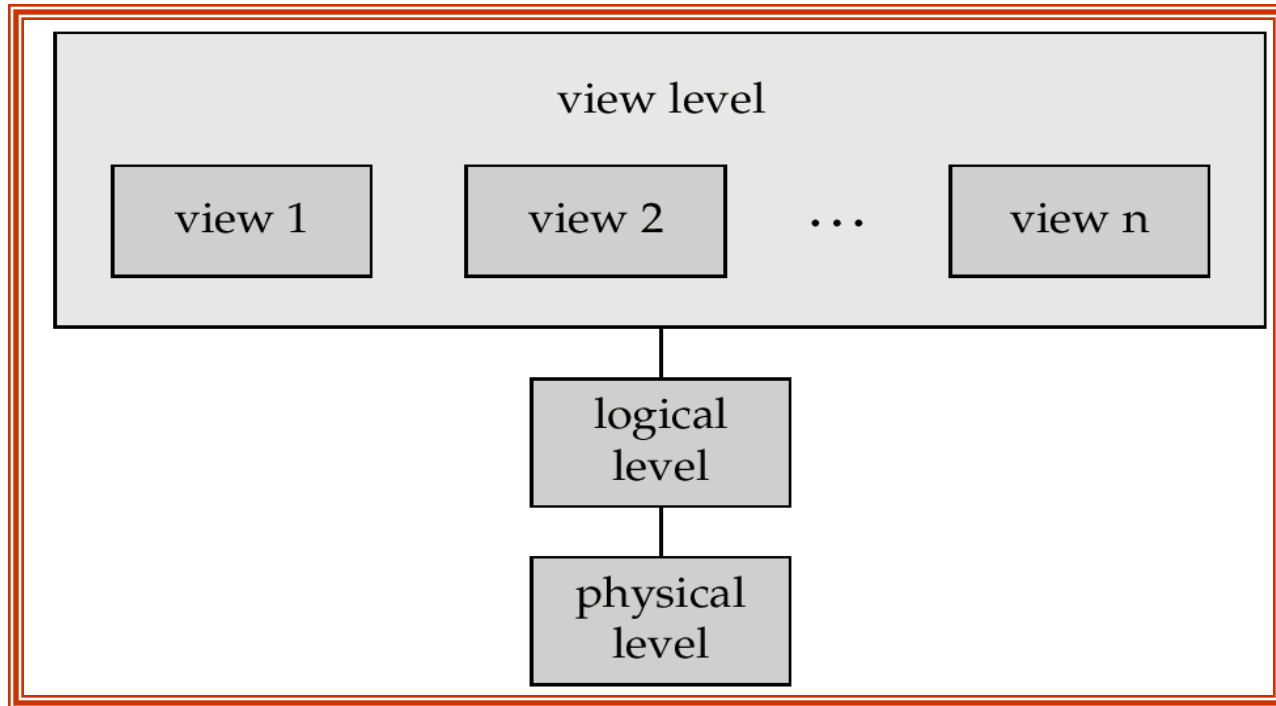


**Key**

This whole table is an **instance** of the Customer schema.

# View of Data

A logical architecture for a database system



Data independence!

# Data retrieval

## ■ Queries

Query = Declarative data retrieval

*describes **what** data to retrieve, not **how** to retrieve it*

*Ex. Give me the students with  $GPA > 3.5$  vs*

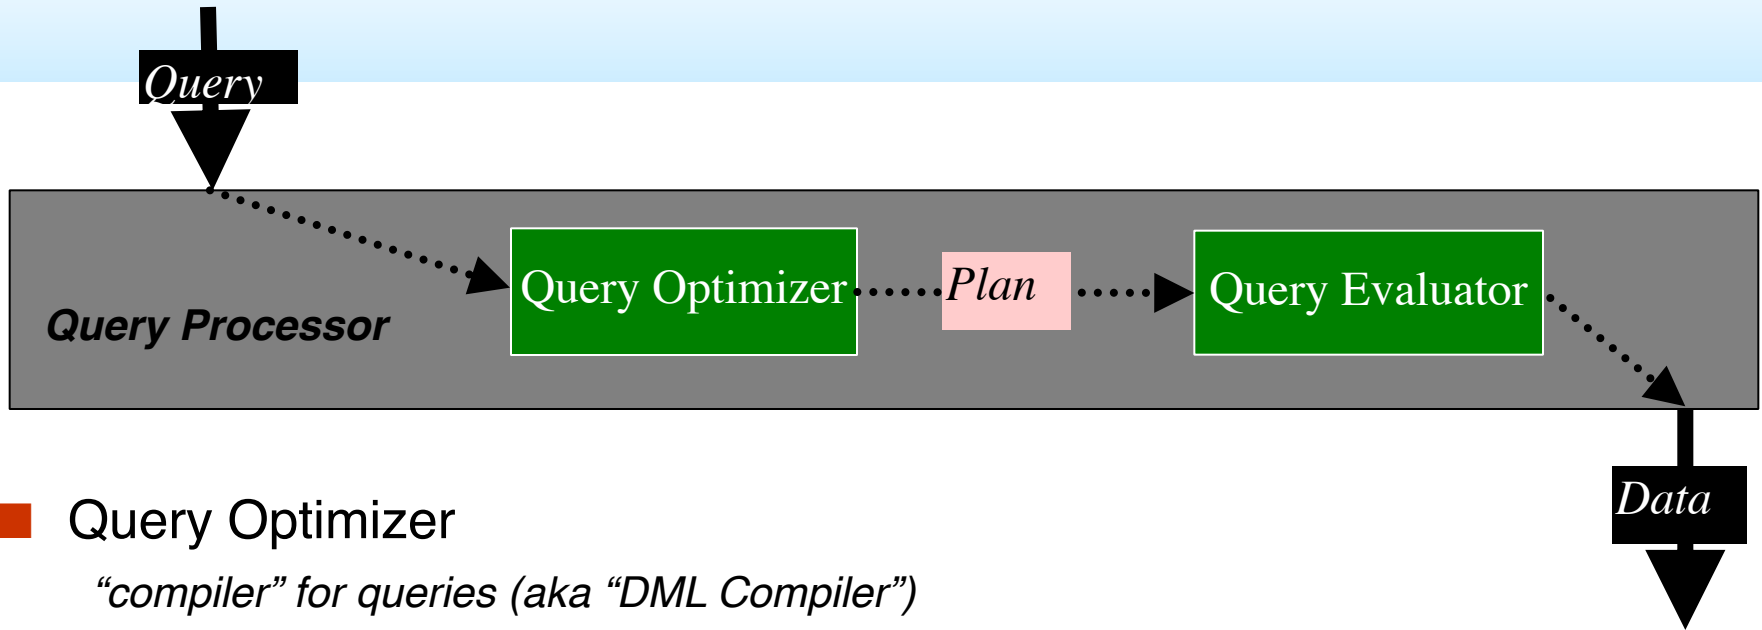
*Scan the student file and retrieve the records with  $GPA > 3.5$*

## ■ Why?

1. Easier to write
2. More efficient to execute (why?)
3. More amenable to change (create/drop an index)



# Data retrieval



## ■ Query Optimizer

*“compiler” for queries (aka “DML Compiler”)*

*Query Plan similar to Assembly Language Program*

Optimizer can do better With declarative queries:

1. *Algorithmic Query (e.g., in C)  $\Rightarrow$  1 Plan to choose from*
2. *Declarative Query (e.g., in SQL)  $\Rightarrow$  n Plans to choose from*

# SQL

## ■ SQL: widely used (declarative) non-procedural language

↗ E.g. find the name of the customer with customer-id 192-83-7465

```
select  customer.customer-name  
from    customer  
where   customer.customer-id = '192-83-7465'
```

↗ E.g. find the balances of all accounts held by the customer with customer-id 192-83-7465


```
select  account.balance  
from    depositor, account  
where   depositor.customer-id = '192-83-7465' and  
         depositor.account-number = account.account-number
```

# Data Integrity

## *Transaction processing*

- Why must we manage concurrent access to data?

John and Jane withdraw \$50 and \$100 from a common account...



Initial balance is \$300. Final balance...

It depends...

# Data Integrity

## Recovery

Transfer \$50 from account A (\$100) to account B (\$200)

1. A.getBalance()
  2. If(balance<sub>A</sub> > \$50)
  3. balance<sub>A</sub> = balance<sub>A</sub> - 50
  4. A.updateBalance(balance<sub>A</sub> )
  5. B.getBalance()
  6. balance<sub>B</sub> = balance<sub>B</sub> + 50
  7. B.updateBalance(balance<sub>B</sub> )
- ← System crashes....

Recovery management

# Database Administrator (DBA)

- Database design
  - Balance needs of multiple applications
- Database tuning
  - Make applications run faster
  - Pick physical structures
- Manage reliability, correctness, security, ...


# Alternative Approaches

- NoSQL (e.g., Mongo DB, Cassandra, CouchBase)
- Hadoop
- Etc.
  
- They give up some of the things that we just discussed in order to
  - ↗ scale better
  - ↗ be easier to deploy
  
- OPINION: Misguided attitude that it's only code -> thus I'll do it myself.

# The Hardware Matters

- Legacy of disk-based processing
  - ↗ The disk is always the bottleneck
  - ↗ Minimize disk accesses
- These days main memory is getting cheap and plentiful.
  - ↗ Store all data in main memory.
  - ↗ Still need some way to persist data (e.g., disk, NVM)
- For distributed DBMS's, the network is the bottleneck
  - ↗ Minimize message passing and data copying.
- Very fast networks are starting to be affordable.

# Assumption

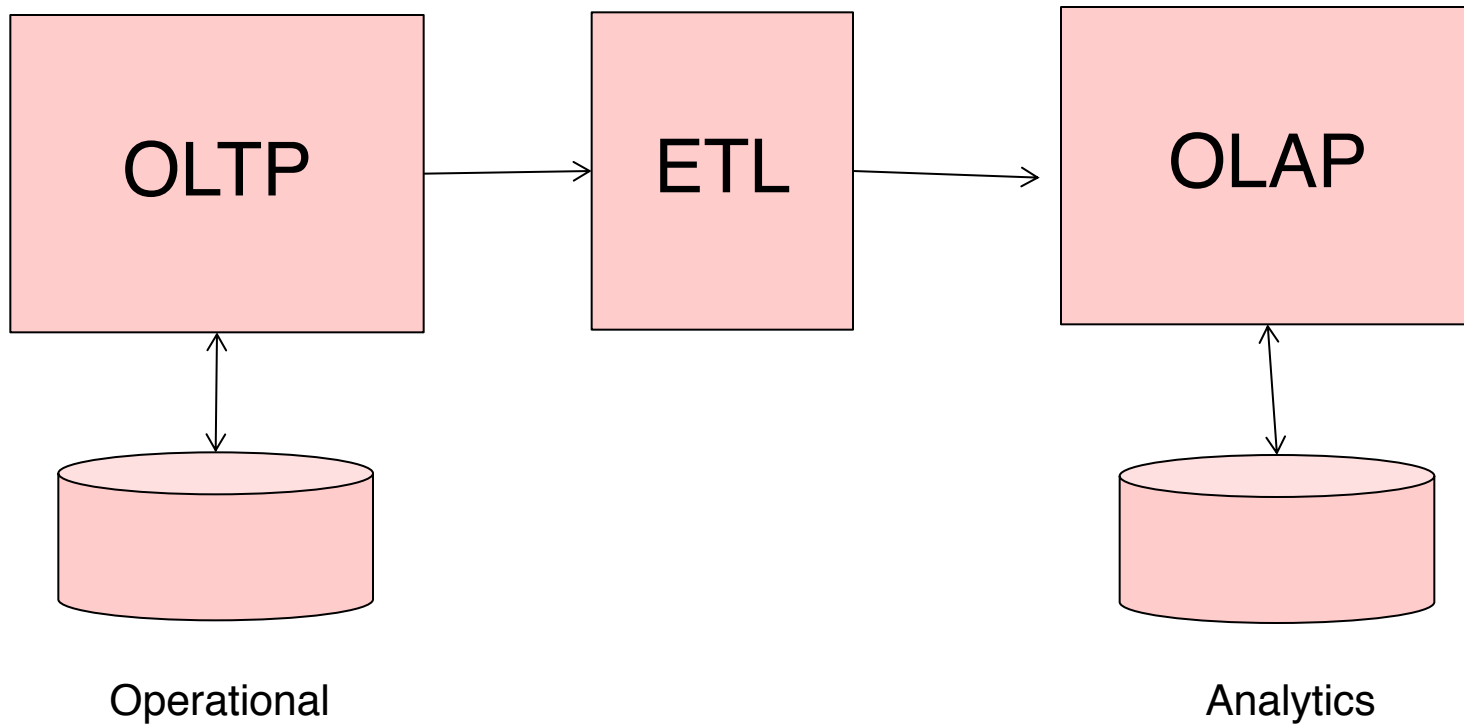





# The Workload Matters

- Workload = the queries that are presented to the system.
- Applications have different characteristics that effect the workload.
- One-Size does not fit all.
  - ↗ OLTP
  - ↗ OLAP
  - ↗ Science

# Reality



# A Closer Look ETL



# Outline

- 1<sup>st</sup> half of the course: application-oriented
  - ↗ How to develop database applications: User + DBA
- 2<sup>nd</sup> part of the course: system-oriented
  - ↗ Learn the internals of a relational DBMS (developer for Oracle..)