

# Halfspaces and Perceptron

## Lecture 2

# Last Time

- We focus on supervised machine learning, with some unsupervised in April
- We use ***empirical risk minimization (ERM)***
  - ERM = pick a hypothesis that minimizes the loss (i.e. empirical risk) on a set of training data
- Naively applying ERM can lead to the pitfall of ***overfitting***
  - Overfitting = picking a hypothesis that is great on training data but very bad on new test data
- Textbook: chapter 1, sections 2.0, 2.1, 2.2

# This Class

- What is a practically useful class of hypotheses that often avoids overfitting?
- How to select an ERM hypothesis from that class computationally efficiently?
- Textbook: sections 9.0, 9.1.0, 9.1.2

Example Revisited:  
Does this animal have cute babies?

# Example Training Data




Tiger

$$\mathbf{z}_1 = (2, 4, 0, 1)$$



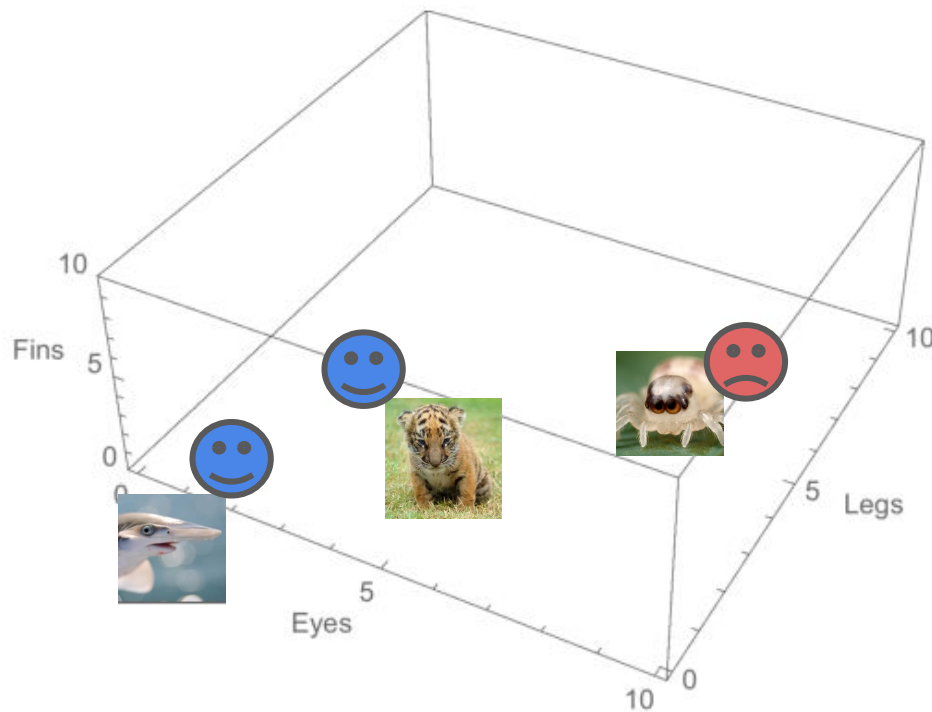
Spider

$$\mathbf{z}_2 = (8, 8, 0, -1)$$



Shark

$$\mathbf{z}_3 = (2, 0, 2, 1)$$

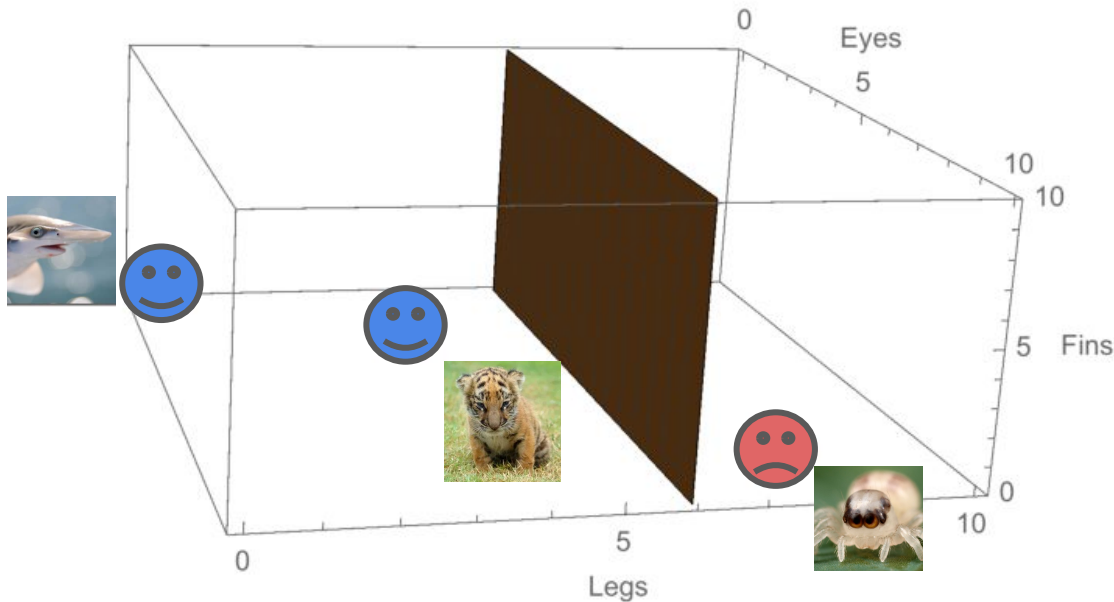


# Example hypothesis


$h = \text{If legs} \leq 6, \text{ then cute}$   
Otherwise, not cute

Equivalent to:

$\text{If } -1 * \text{legs} + 6 > 0, \text{ then cute}$   
Otherwise, not cute

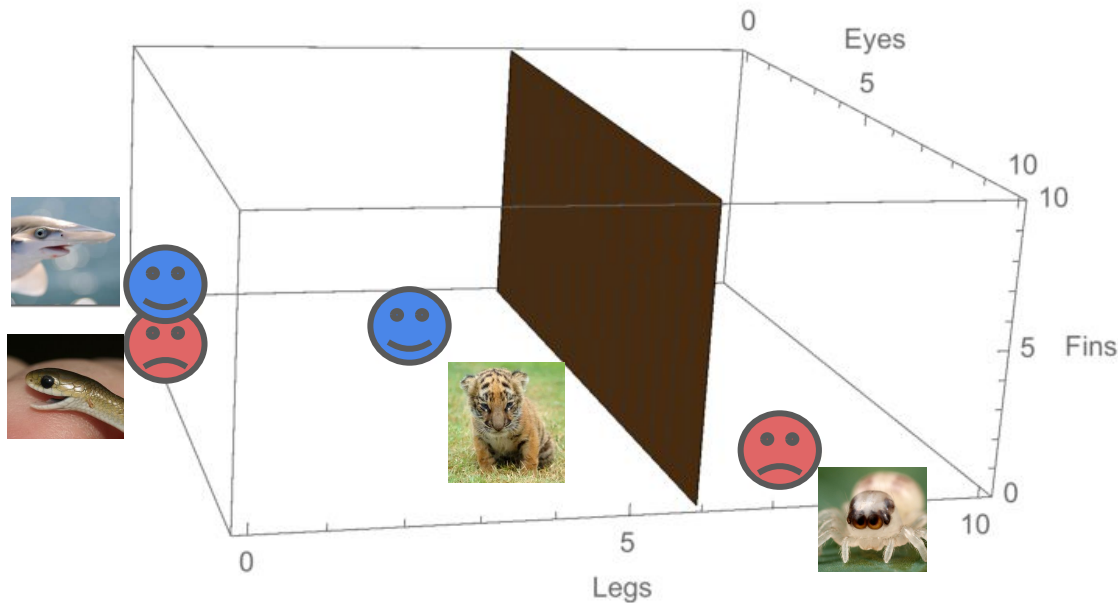


# What if we also had Snake?



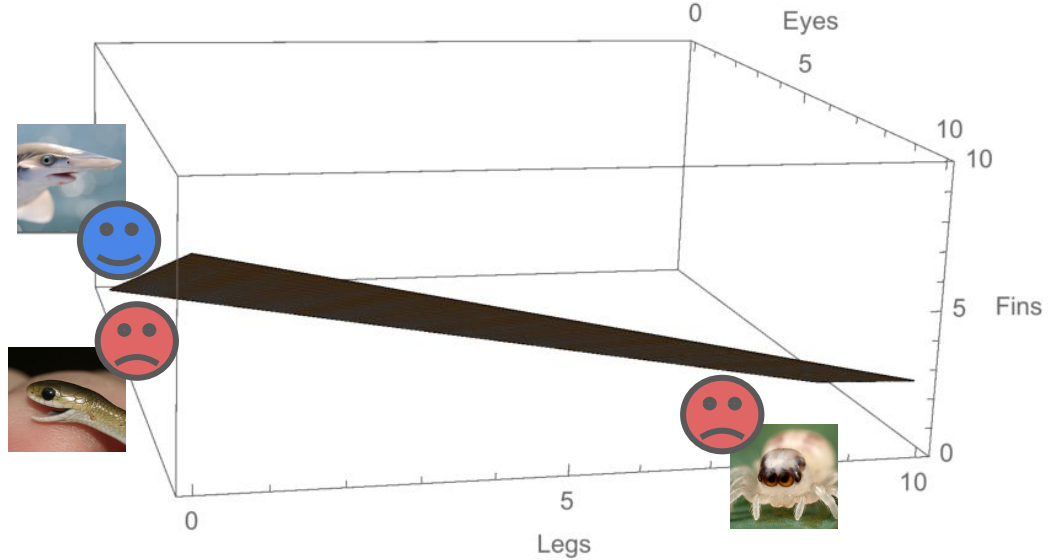
Snake

$$z_4 = (2, 0, 0, -1)$$



# Example improved hypothesis

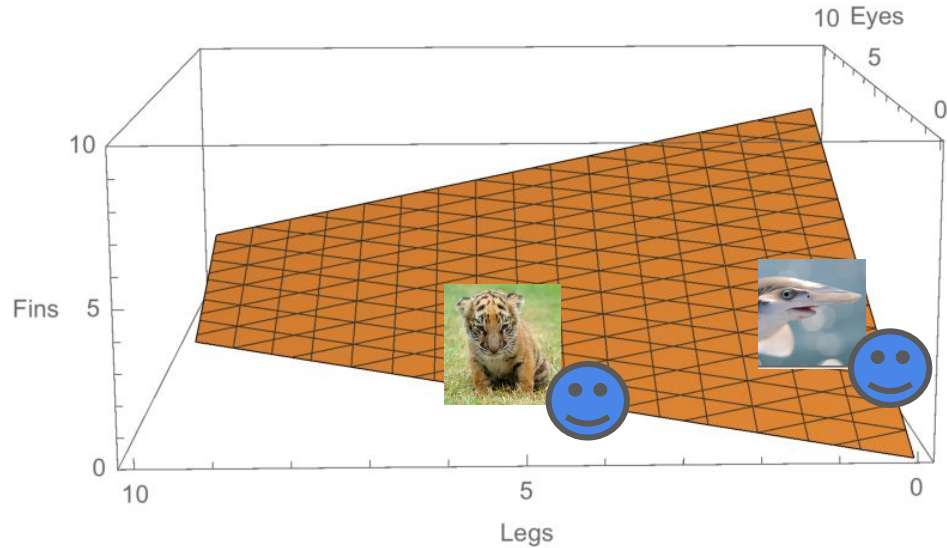
$h = \text{If } -1.5 * \text{Eyes} + \text{Legs} + 2 * \text{Fins} > 0, \text{ then cute}$   
Otherwise, not cute





# Alternate View (Rotated Horizontally 180°)

$h = \text{If } -1.5 * \text{Eyes} + \text{Legs} + 2 * \text{Fins} \geq 0, \text{ then cute}$   
Otherwise, not cute



# Halfspaces

# Affine Function

Fancy name for a linear function plus a constant:

$$h_{\mathbf{w},b}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b$$

where  $\mathbf{x} \in \mathbb{R}^d, b \in \mathbb{R}$

Often, we will omit  $b$  and assume that the last element of each  $\mathbf{x}$  is 1

# Halfspace Hypothesis

We can use an affine function to define a hypothesis called a halfspace:

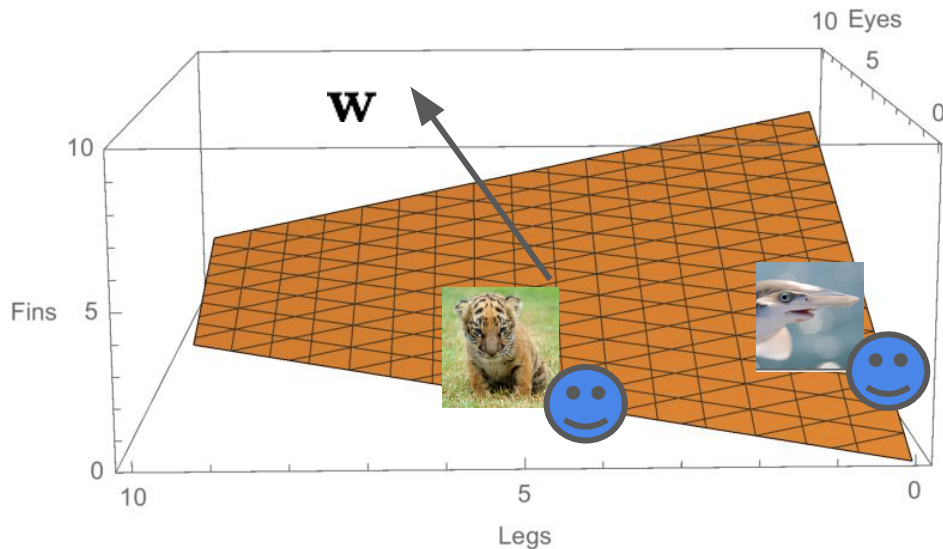
$$h_{\mathbf{w}}(\mathbf{x}) = \text{sign}(\langle \mathbf{w}, \mathbf{x} \rangle)$$

$$\mathcal{X} = \mathbb{R}^d \quad \mathcal{Y} = \{1, -1\}$$

# Decision Boundary

In a  $d$ -dimensional attribute space  
the decision boundary is the  
 $(d-1)$ -dimensional hyperplane  
where  $\langle \mathbf{w}, \mathbf{x} \rangle = 0$

$\mathbf{w}$  is a vector normal to that  
hyperplane (pointing towards  
the positive side)



# Question



# How We'll Do Questions: Peer Instruction

1. Solo Vote
2. Group Discussion
3. Group Vote
4. Class Discussion

***Remember:*** You get full credit for participating, not selecting the correct answer

And backed by SCIENCE!

[A Multi-institutional Study of Peer Instruction in Introductory Computing](#)

Leo Porter, Dennis Bouvier, Quintin Cutts, Scott Grissom, Cynthia Lee, Robert McCartney, Daniel Zingaro, Beth Simon.  
*SIGCSE 2016, March 2016.*

We'll Have to Wait for the Question





# Linear Separation and Realizability

# Linear Separation

- What types of training data can we separate perfectly with a halfspace?
- I.e., when does there exist  $\mathbf{w}$  such that  $L_{\mathcal{D}}(h_{\mathbf{w}}) = 0$ ?
- More generally, beyond halfspace classifiers, called the **realizability assumption**
- Note that the realizability assumption implies  $L_{\mathcal{S}}(h_{\mathbf{w}}) = 0$  for all  $S \sim \mathcal{D}$


# Example

Let  $f : \mathcal{X} \rightarrow \mathcal{Y}$  be the true labeling function, i.e.,  $\mathcal{D}(y = f(\mathbf{x})|\mathbf{x}) = 1$   
(and let  $\mathcal{D}(\mathbf{x})$  be uniform)

Suppose  $\mathcal{X} = \{0, 1\}^d$ ,  $\mathcal{Y} = \{1, -1\}$ ,  
and

$$f(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_{i=1}^d x_i \geq 1 \\ -1 & \text{otherwise} \end{cases}$$

Linearly Separable!



# Question



We'll Have to Wait for the Question



# Perceptron Algorithm

# Intuition

- Iterative algorithm
- Initialize weights to be all zeroes
- For each misclassified point, add the point (signed by its label) to the weight vector to shift it
- Continue until all points are classified correctly (or stuck)

# Perceptron Training Algorithm

```
 $w^{(1)} = (0, 0, \dots, 0)$   
for  $t = 1, 2, \dots$   
  if  $(\exists i \text{ s.t. } y_i \langle w, x_i \rangle \leq 0)$   
     $w^{(t+1)} = w^{(t)} + y_i x_i$   
  else  
    return  $w^{(t)}$ 
```




# Animation Example

<https://www.youtube.com/watch?v=xpJHhHwR4DQ>

# Why Does it Work?

- Halfspace is correct if  $y\langle \mathbf{w}, \mathbf{x} \rangle > 0$
- Suppose for some training example that  $y\langle \mathbf{w}, \mathbf{x} \rangle \leq 0$
- Adding  $y\mathbf{x}$  to  $\mathbf{w}$  moves weights toward right direction
- Always helps:

$$\begin{aligned} y\langle \mathbf{w}^{(t+1)}, \mathbf{x} \rangle &= y\langle \mathbf{w}^{(t)} + y\mathbf{x}, \mathbf{x} \rangle \\ &= y\langle \mathbf{w}^{(t)}, \mathbf{x} \rangle + \|\mathbf{x}\|_2^2 \end{aligned}$$



# Formal Guarantee

**THEOREM 9.1** *Assume that  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$  is separable, let  $B = \min\{\|\mathbf{w}\| : \forall i \in [m], y_i \langle \mathbf{w}, \mathbf{x}_i \rangle \geq 1\}$ , and let  $R = \max_i \|\mathbf{x}_i\|$ . Then, the Perceptron algorithm stops after at most  $(RB)^2$  iterations, and when it stops it holds that  $\forall i \in [m], y_i \langle \mathbf{w}^{(t)}, \mathbf{x}_i \rangle > 0$ .*


Proof sketch: let  $\mathbf{w}^*$  be the weight vector that achieves the min. in the definition of  $B$  and show that after  $T$  iterations:

$$\frac{\langle \mathbf{w}^*, \mathbf{w}^{(T+1)} \rangle}{\|\mathbf{w}^*\| \|\mathbf{w}^{(T+1)}\|} \geq \frac{\sqrt{T}}{RB}$$

Do Halfspaces Ever Overfit?


# Halfspaces seem pretty great, right?

- If the data is linearly separable, then it's awesome
- If it's not, we can get something like this:



# More Dimensions can Help

- Think back to the xor labeling function:  $f^C(\mathbf{x}) = x_1 \text{ xor } x_2$
- We could make such a 2-d data set into a 3-d linearly separable data set through feature engineering, i.e., expanding  $\mathcal{H}$  to operate on new attributes  $\mathbf{x}' = \psi(\mathbf{x})$  where  $\psi(\mathbf{x}) = (x_1, x_2, x_1 \cdot x_2)$



# Too Many Dimensions can Lead to Overfitting

- We can place  $d+1$  points in  $\mathbb{R}^d$  and linearly separate **any** labels assigned to them
- Example: if our points are  $\mathbf{0}, \mathbf{e}_1 = (1, 0, 0, \dots), \mathbf{e}_2 = (0, 1, 0, \dots), \dots$  then for any  $y_1, y_2, \dots$ , set  $b$  to  $y_1$  and  $\mathbf{w}$  to  $y_2, y_3, \dots$
- The technical term for this is ***shattering***

# So How Do We Know What to Do?

- Most of supervised machine learning is balancing the “triple tradeoff” (Dietterich, 2003) of the complexity of  $\mathcal{H}$ , the size of the training data  $m$ , and the true error  $L_{\mathcal{D}}(h_S)$
- Throughout the semester we’ll see formally and in practice that we need domain knowledge to balance this triple tradeoff. There is no free lunch!



# Our first encounter with the curse of dimensionality...

I'll just leave this here: <https://www.youtube.com/watch?v=DQWI1kvmwRg>

# The Most Important Things

- **Halfspaces** are hypotheses defined as hyperplanes that separate two classes
- Training data that can be separated perfectly is **linearly separable**
  - More general term for any hypothesis class with perfect fit for training data: **realizable**
- The **perceptron** algorithm is a simple method for learning halfspaces
  - Finds ERM solution efficiently if training data is linearly separable
- Textbook: sections 9.0, 9.1.0, 9.1.2

# Next Class

- How can we build linear predictors for predicting continuous values?
- Textbook: sections 9.2