

Linear and Polynomial Regression

Lecture 3

Last Time

- **Halfspaces** are hypotheses defined as hyperplanes that separate two classes
- Training data that can be separated perfectly is **linearly separable**
 - More general term for any hypothesis class with perfect fit for training data: **realizable**
- The **perceptron** algorithm is a simple method for learning halfspaces
 - Finds ERM solution efficiently if training data is linearly separable
- Textbook: sections 9.0, 9.1.0, 9.1.2

This Class

- How can we build linear predictors for predicting continuous values?
- Textbook: section 9.2

Linear Regression

Regression

- ***Regression**** refers to modeling a relationship (often one-to-one, usually smooth) between a target variable and other variables
- *A term often used in subtly confusing ways. For example, we'll talk about a regression model called "logistic regression" that predicts the probability of a discrete label, essentially a classification task. Don't sweat it too much.

Continuous Regression

What if we want to learn a program that outputs continuous values?

Last class: $\mathcal{Y} = \{1, -1\}$

Now: $\mathcal{Y} = \mathbb{R}$

What applications can you think of?


Linear Regression

Linear regression is a common hypothesis class for the following domain and label set:

$$\mathcal{X} = \mathbb{R}^d \quad \mathcal{Y} = \mathbb{R}$$


Again we use a linear predictor, but now its value is the output (instead of taking the sign)

$$h_{\mathbf{w},b}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle + b$$




Example: Baby Weight

Weight (lbs.) vs. Month



Example: Baby Weight

Weight (lbs.) vs. Month



● — $1.15 * x + 9.74$

ML algorithm = representation
+ loss function + optimizer

Loss for Linear Regression

- Squared loss, a.k.a. mean squared error (MSE):

$$L_S(h_{\mathbf{w}}) = \frac{1}{m} \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

- Another option is absolute value loss:

$$L_S(h_{\mathbf{w}}) = \frac{1}{m} \sum_{i=1}^m |h_{\mathbf{w}}(\mathbf{x}_i) - y_i|$$

Question



We'll Have to Wait for the Question



ML algorithm = representation
+ loss function + optimizer

Optimizer for Linear Regression

- Least Squares

$$\arg \min_{\mathbf{w}} L_S(h_{\mathbf{w}}) = \arg \min_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

- How do we solve this optimization problem?

Least Squares

- Want to solve

$$\arg \min_{\mathbf{w}} L_S(h_{\mathbf{w}}) = \arg \min_{\mathbf{w}} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$$

- We start by finding the gradient of the objective with respect to \mathbf{w} and setting it equal to the zero vector

$$\frac{2}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i) \mathbf{x}_i = \mathbf{0}$$

- Calculus tells us that this is a necessary condition for any optimizer \mathbf{w}_S

Rewriting the Least Squares Condition

A useful way to rewrite $\frac{2}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i) \mathbf{x}_i = \mathbf{0}$ in equivalent forms:

$$A\mathbf{w} = \mathbf{b} \quad \text{where} \quad A = \left(\sum_{i=1}^m \mathbf{x}_i \mathbf{x}_i^\top \right) \quad \text{and} \quad \mathbf{b} = \sum_{i=1}^m y_i \mathbf{x}_i$$

or

$$A = \begin{pmatrix} \vdots & & \vdots \\ \mathbf{x}_1 & \dots & \mathbf{x}_m \\ \vdots & & \vdots \end{pmatrix} \begin{pmatrix} \vdots & & \vdots \\ \mathbf{x}_1 & \dots & \mathbf{x}_m \\ \vdots & & \vdots \end{pmatrix}^T \quad \text{and} \quad \mathbf{b} = \begin{pmatrix} \vdots & & \vdots \\ \mathbf{x}_1 & \dots & \mathbf{x}_m \\ \vdots & & \vdots \end{pmatrix} \begin{pmatrix} y_1 \\ \vdots \\ y_m \end{pmatrix}$$

If A is invertible...

- Easy! One unique ERM: $\mathbf{w}_S = A^{-1}\mathbf{b}$
- A is invertible if and only if $\mathbf{x}_1, \dots, \mathbf{x}_m$ span \mathbb{R}^d .

If A is not invertible...

It's symmetric, so we can do an eigenvalue decomposition: $A = VDV^{\top}$

Now define D^{+} where D_{ij}^{+} is D_{ij}^{-1} if D_{ij} is nonzero and 0 otherwise

Next define $A^{+} = VD^{+}V^{\top}$ (Moore-Penrose Inverse)

Then

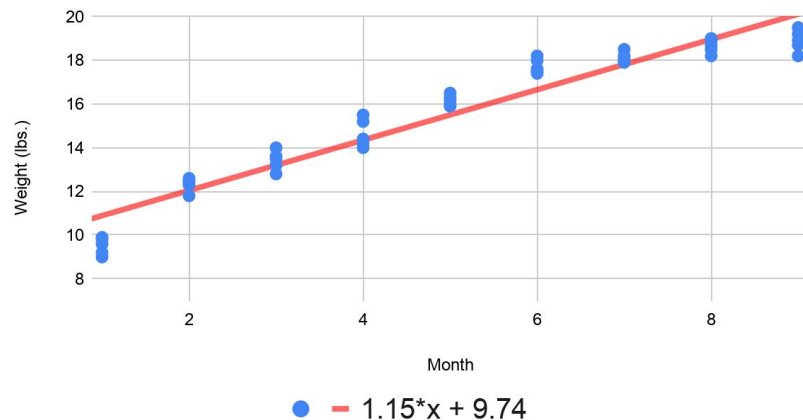
$$\mathbf{w}_S = A^{+}\mathbf{b}$$

Polynomial Regression


Polynomial Regression

- By expanding the representation, and therefore the hypothesis class, can also fit a more general polynomial function of the original features

Weight (lbs.) vs. Month



Weight (lbs.) vs. Month




Polynomial Regression

- Just like we made xor linearly separable, we can reduce empirical risk by transforming $\mathbf{x}' = \psi(\mathbf{x})$ and running the same algorithm
- Let's choose $\psi(\mathbf{x})$ to be the original features raised to increasing powers, i.e., if $\mathcal{X} = \mathbb{R}$ then $\psi(x) = (x, x^2, x^3)$
- Then running linear regression on $x' = \psi(x)$ is the same as learning the weights of a polynomial of this form:

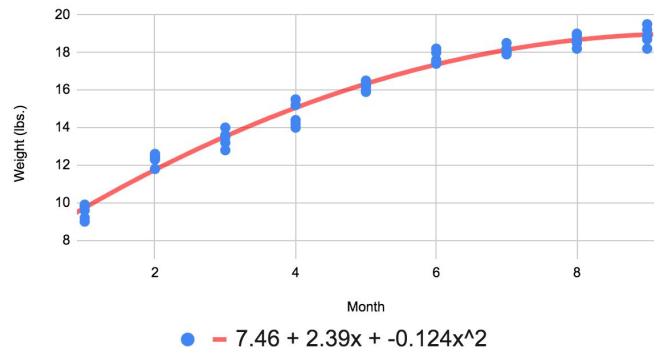
$$h(x) = w_1 + w_2x + w_3x^2 + w_4x^3 + \cdots + w_{p+1}x^p$$

Polynomial Regression Demo


Weight (lbs.) vs. Month




Weight (lbs.) vs. Month



Weight (lbs.) vs. Month



Weight (lbs.) vs. Month



Question



We'll Have to Wait for the Question



PSA from your Friendly Neighborhood Statisticians

In machine learning, modeling assumptions are rarely causal

Use extreme caution when interpreting a hypothesis or predicting counterfactuals!

The Most Important Things

- ***Linear regression*** is a hypothesis class for predicting a continuous value from a linear combination of the attributes
- ***Polynomial regression*** generalizes linear regression. Same loss and optimizer, but we expand representation by taking polynomials of chosen degree of each attribute (still linear in weights so same algorithm!)
- Textbook: section 9.2

Next Class

- How can we build linear predictors
 - for predicting probabilities of discrete classes?
 - for predicting more than two classes?
- Textbook: Chapters 9.3 12.1.1, 14.0, 14.1.0